



euricom

A DIMENSION DATA COMPANY

accelerate
your ambition

Containers and Docker



Lab 1

001-Lab-Setup

Lab Goals:


- Set up infrastructure
- Get acquainted with Cloud Shell / Minikube
- Explore the git repo

Before Containers

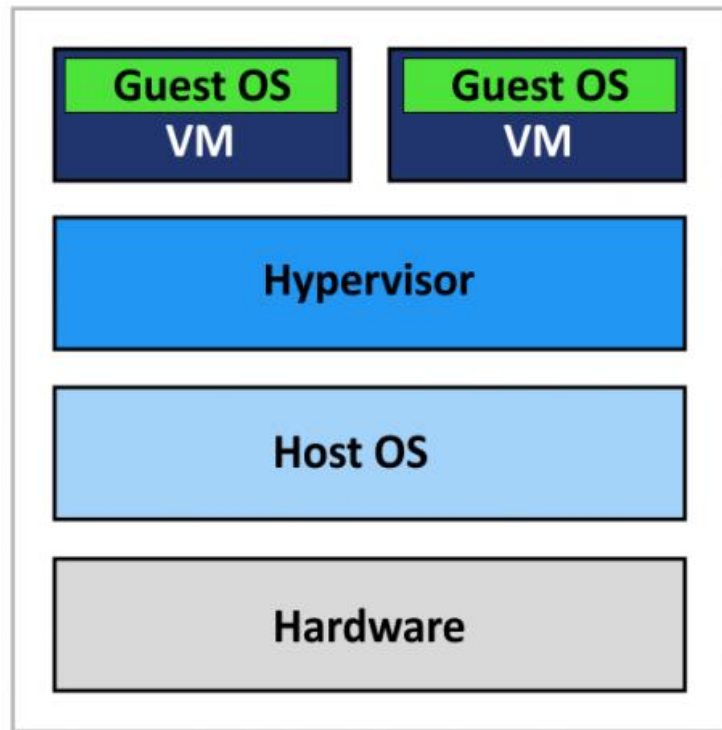
- 'Big' and 'fast' servers
- 1 server = 1 application
- Migration?
- Cost?
- Environment Replication?
- Scalable?

Virtual Machines (VMs)

- 1 server = multiple applications
- Hypervisor (VM Monitor) Architecture
 - Type 1 (Windows Hyper-V)
 - Type 2 (VMWare)
- Host hardware allocation across all apps
- Each VM contains:
 - Hardware (Virtual)
 - OS
 - Application

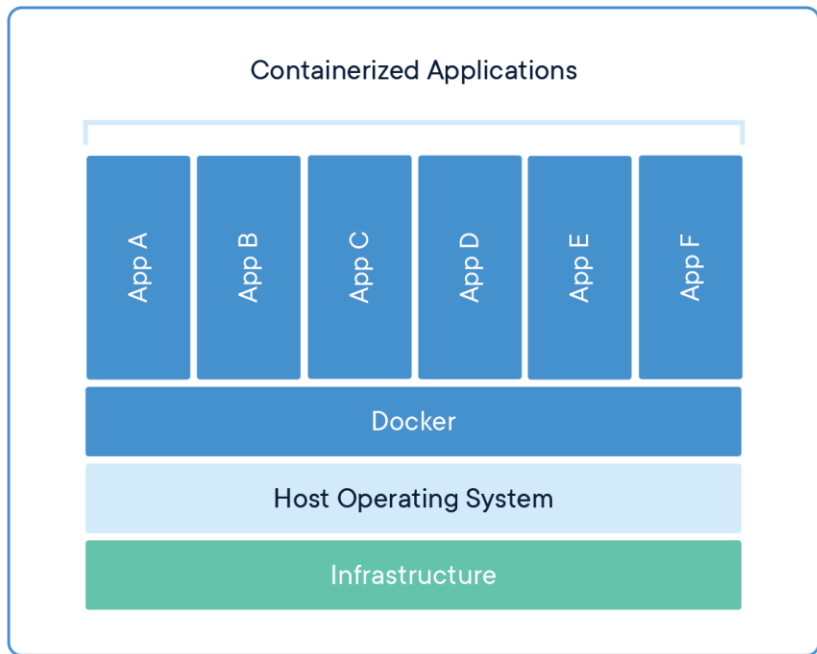


**Type 1 Hypervisor
(Bare-Metal Architecture)**



**Type 2 Hypervisor
(Hosted Architecture)**

Containers



- Standardised unit of software
 - Application layer construction
- Own libraries and packages
- Shared OS

Containers

- Standardisation and productivity
- Compatibility and maintainability
- CD/CI

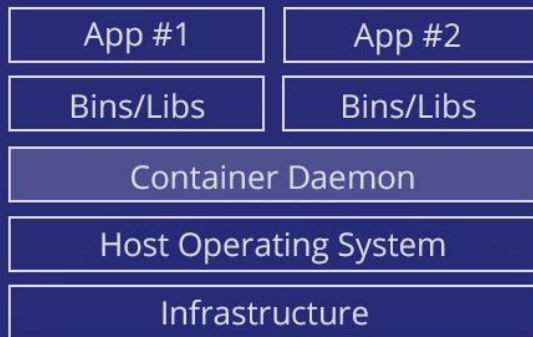
VMs VS Containers

VIRTUAL MACHINES



WHATS
—*the*—
DIFF?

CONTAINERS



Docker 101



Image

The basis of a Docker container. The content at rest.



Container

The image when it is 'running.' The standard unit for app service



Engine

The software that executes commands for containers. Networking and volumes are part of Engine. Can be clustered together.



Registry

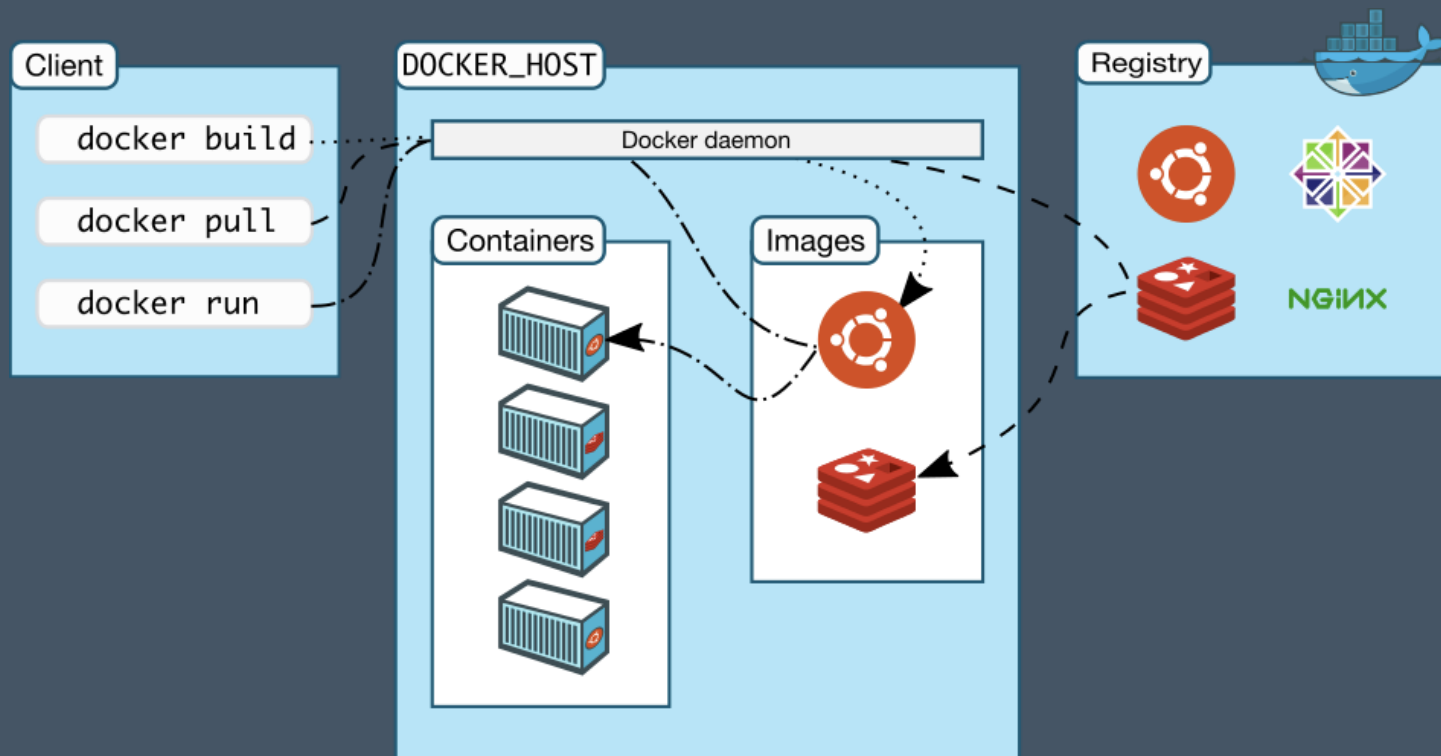
Stores, distributes and manages Docker images



Control Plane


Management plane for container and cluster orchestration

Docker Engine



DockerFile

- Text document for building docker image
- Contains all CLI instructions for the build





Lab 2

002-Containerizing-An-Application

Lab Goals:

- Run Golang API locally
- Build Docker Image
- Run the Docker Container
- Use Environment Variables



euricom


A DIMENSION DATA COMPANY

accelerate
your ambition

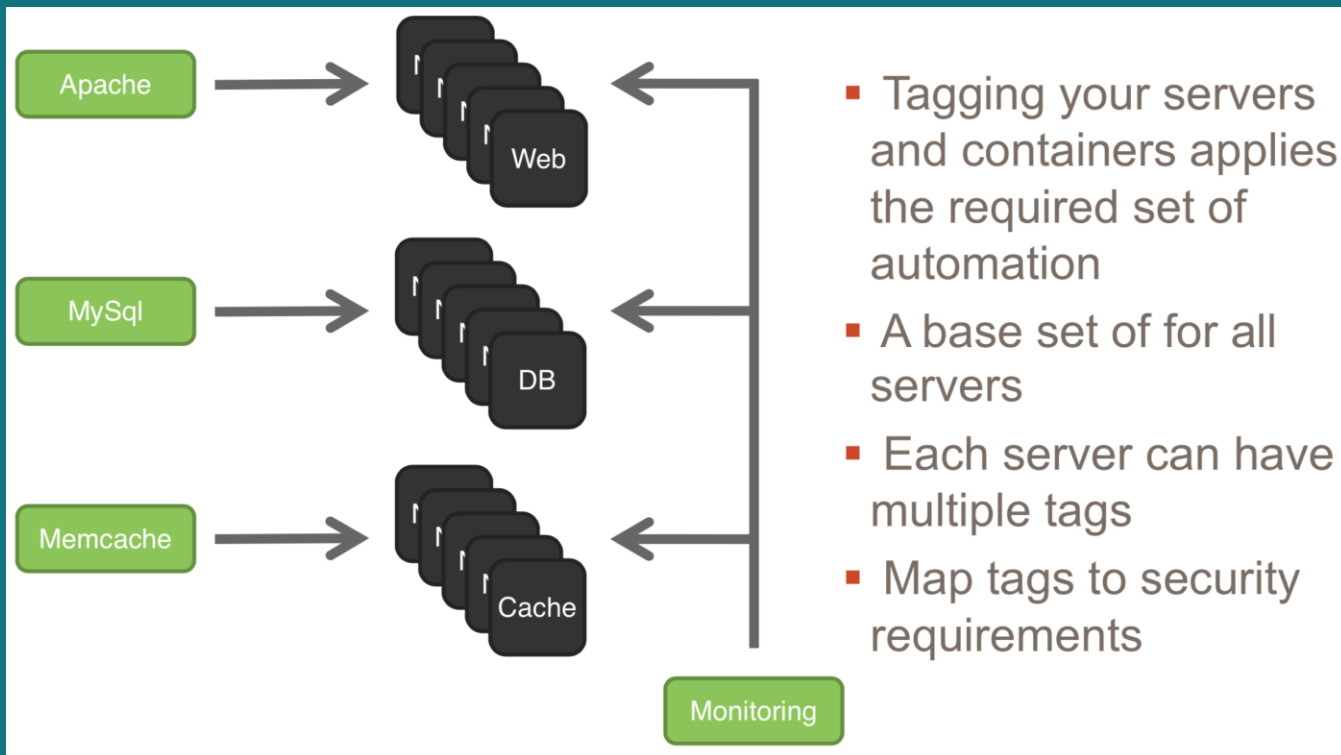
An introduction to Kubernetes (k8s)

Infrastructure as Code (intro to k8s)

Infrastructure as code (IaC) allows for infrastructure to be deployed using a high-level, descriptive language. IaC treats the entire infrastructure as if it is software. Because, it's all software...



Grouping and Tagging



Cattle, not pets



Security wins !

- Security team now has insight into the entire system
- Infrastructure is auditable and version controlled, just like source code
- Patching can be applied programmatically with a high level of certainty
- Alerting can be built for changes to specific areas of the infrastructure
- A new firewall rule is created or deleted – Administrative user is created
 - New VPC rolled out
- Testing can occur much earlier in the pipeline

Infrastructure as Code - Terraform



Learn how Terraform fits into the HashiCorp Suite

[Intro](#)[Learn](#)[Docs](#) [Community](#)[Enterprise](#)[Download](#)[GitHub](#)[Sign In](#)

NEW

[Introducing Terraform Cloud](#) →

Terraform

Use Infrastructure as Code to provision and manage any cloud, infrastructure, or service

[Sign up for Cloud](#)[Download CLI](#)

```
Last login: Wed Sep  4 11:20:02 on ttys000
terraformUser$
```

Infrastructure as Code – K8s

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-site-ingress
  namespace: my-site-prod
  annotations:
    kubernetes.io/tls-acme: "true"
    kubernetes.io/ingress.class: "gce"
    kubernetes.io/ingress.global-static-ip-name: my-site-external-ip
spec:
  tls:
  - hosts:
    - api.my.site
    - my.site
    secretName: my-site-cert
  rules:
  - host: api.my.site
    http:
      paths:
      - path: /*
        backend:
          serviceName: app-api
          servicePort: 80
  - host: my.site
    http:
      paths:
      - path: /*
        backend:
          serviceName: my-site-prod
          servicePort: 80
```

Kubernetes is an open-source platform built to automate **deployment, scaling and orchestration** of containers.

K8S is **portable**. Clusters can be deployed on a public/private cloud, on prem, and even on your laptop.

K8S is **customizable**. It is modular and extensible to fit a variety of use-cases.


K8S is **scalable**. It provides self-healing, auto scaling, and replication out of the box.

**virtual
machines that
Kubernetes
manages**

cluster



cluster




cluster



pod

**group of
containers
sharing storage
and network**


pod




pod

```
apiVersion: v1
kind: Pod
metadata:
  name: redis-rails
spec:
  containers:
  - name: key-value
    image: redis
    ports:
    - containerPort: 6379
  - name: rails-frontend
    image: rails
    ports:
    - containerPort: 3000
```

pod.yaml



pod.yaml



pod.yaml

master

node

node



node

pod.yaml



deployment

**ensure N pods
are up and
running**

deployment

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: rails-deployment
  labels:
    app: rails
spec:
  replicas: 4
  selector:
    matchLabels:
      app: rails
  template:
    metadata:
      labels:
        app: rails
    spec:
      containers:
        - name: key-value
          image: redis
          ports:
            - containerPort: 6379
        - name: rails-frontend
          image: rails
          ports:
            - containerPort: 3000
```

deploy.yaml

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: rails-deployment
  labels:
    app: rails
spec:
  replicas: 4
  selector:
    matchLabels:
      app: rails
  template:
    metadata:
      labels:
        app: rails
    spec:
      containers:
        - name: key-value
          image: redis
          ports:
            - containerPort: 6379
        - name: rails-frontend
          image: rails
          ports:
            - containerPort: 3000
```

deploy.yaml

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: rails-deployment
  labels:
    app: rails
spec:
  replicas: 4
  selector:
    matchLabels:
      app: rails
  template:
    metadata:
      labels:
        app: rails
    spec:
      containers:
        - name: key-value
          image: redis
          ports:
            - containerPort: 6379
        - name: rails-frontend
          image: rails
          ports:
            - containerPort: 3000
```


deploy.yaml

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: rails-deployment
  labels:
    app: rails
spec:
  replicas: 4
  selector:
    matchLabels:
      app: rails
  template:
    metadata:
      labels:
        app: rails
    spec:
      containers:
        - name: key-value
          image: redis
          ports:
            - containerPort: 6379
        - name: rails-frontend
          image: rails
          ports:
            - containerPort: 3000
```


deploy.yaml

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: rails-deployment
  labels:
    app: rails
spec:
  replicas: 4
  selector:
    matchLabels:
      app: rails
  template:
    metadata:
      labels:
        app: rails
    spec:
      containers:
        - name: key-value
          image: redis
          ports:
            - containerPort: 6379
        - name: rails-frontend
          image: rails
          ports:
            - containerPort: 3000
```


deploy.yaml




deploy.yaml



deploy.yaml



10.0.0.1



10.0.0.2

10.0.0.3



10.0.0.4

deploy.yaml




service

**abstraction
layer that
enables pod
communication**

service




10.0.0.1




10.0.0.2

10.0.0.3




10.0.0.4

service



service



service

your.site.com

Public load balancer

Service

master


node

node

node

Service

service



```
kind: Service
apiVersion: v1
metadata:
  name: web-frontend
spec:
  ports:
    - name: http
      port: 80
      targetPort: 3000
      protocol: TCP
  selector:
    app: rails
  type: LoadBalancer
```

svc.yaml

```
kind: Service
apiVersion: v1
metadata:
  name: web-frontend
spec:
  ports:
    - name: http
      port: 80
      targetPort: 3000
      protocol: TCP
  selector:
    app: rails
type: LoadBalancer
```

svc.yaml

```
kind: Service
apiVersion: v1
metadata:
  name: web-frontend
spec:
  ports:
    - name: http
      port: 80
      targetPort: 3000
      protocol: TCP
  selector:
    app: rails
  type: LoadBalancer
```

svc.yaml



Labels and Selectors

**Metadata (key-value) which
can be attached
to a resource**

Labels

**Used for
identification
such as app
name, tier,
environment**

Labels

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: rails-deployment
  labels:
    app: rails
spec:
  replicas: 4
  selector:
    matchLabels:
      app: rails
  template:
    metadata:
      labels:
        app: rails
    spec:
      containers:
        - name: key-value
          image: redis
          ports:
            - containerPort: 6379
        - name: rails-frontend
          image: rails
          ports:
            - containerPort: 3000
```

deploy.yaml

**Provides loose
coupling**

between objects

Selectors

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: rails-deployment
  labels:
    app: rails
spec:
  replicas: 4
  selector:
    matchLabels:
      app: rails
  template:
    metadata:
      labels:
        app: rails
    spec:
      containers:
        - name: key-value
          image: redis
          ports:
            - containerPort: 6379
        - name: rails-frontend
          image: rails
          ports:
            - containerPort: 3000
```

deploy.yaml



Ingress

**configure
external access
to your cluster**

Ingress

```
kind: Ingress
apiVersion: extensions/v1beta1
metadata:
  name: web-ingress
spec:
  backend:
    serviceName: web-frontend
    servicePort: 80
```

ingress.yaml

```
kind: Ingress
apiVersion: extensions/v1beta1
metadata:
  name: web-ingress-vhosts
  rules:
    - host: sub.domain.com
      http: paths:
        - backend:
            serviceName: web-frontend-1
            servicePort: 80
    - host: other.domain.com
      http: paths:
        - backend:
            serviceName: web-frontend-2
            servicePort: 80
```

ingress.yaml

**manage
different
environments in
the same
cluster**

namespace

```
kind: Namespace
apiVersion: v1
metadata:
  name: development
```

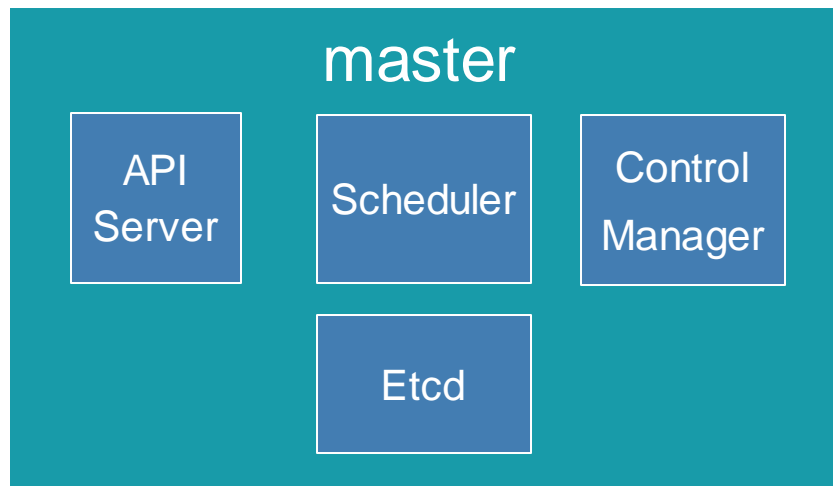
ns.yaml



K8s internals



cluster



master

master

API
Server

master

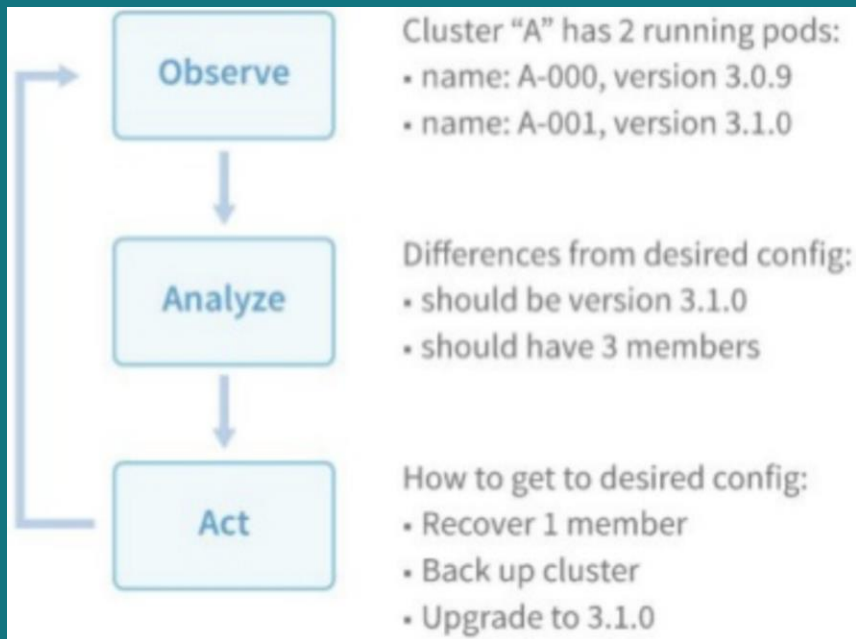
master

Control
Manager

master

Controllers

Perform reconciliation loops on cluster resources and converge *desired* state with the *actual* state.



master

Control Manager

Node Controller

Replication Controller

ServiceAccount Controller

...

master

Node Controller

Responsible for noticing when nodes go up and down

Replication Controller

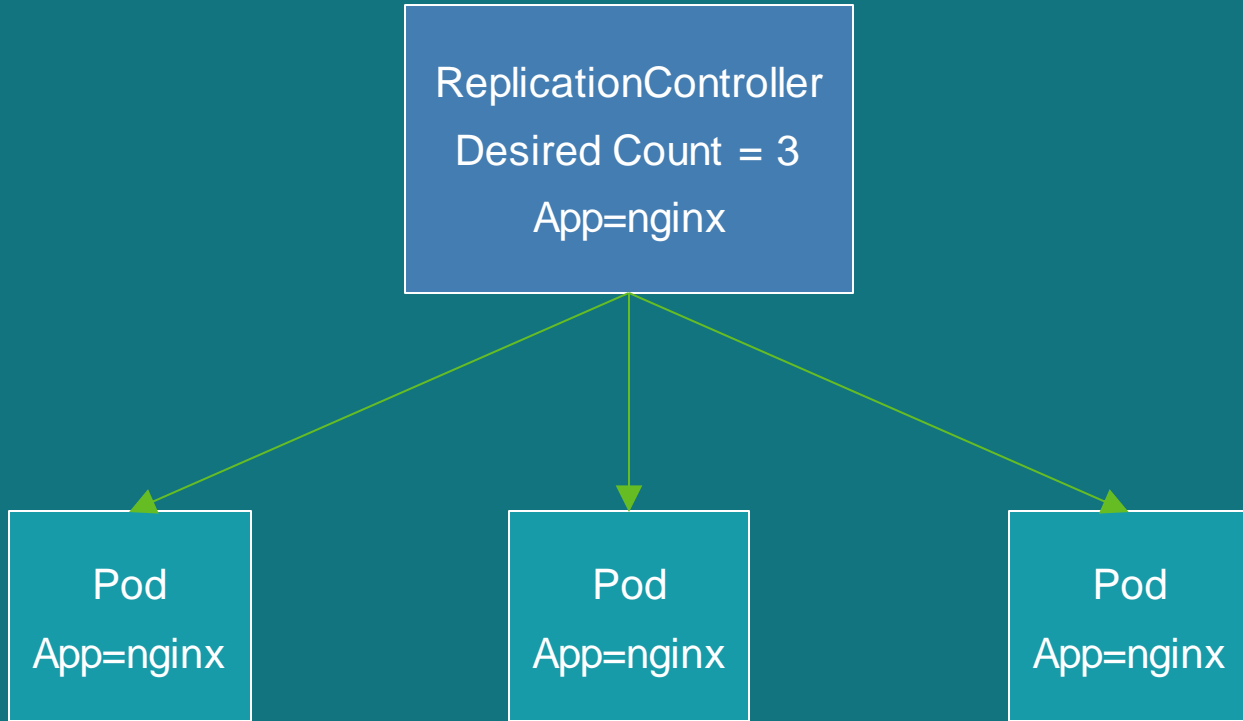
Responsible for maintaining the correct number of pods for every replication controller object

ServiceAccount Controller

Creates default accounts and API access tokens

...


master



master

Etcd

master



node

node

kube-proxy

kubelet

Container runtime

node

node

kube-proxy

node

node

kubelet

node

node

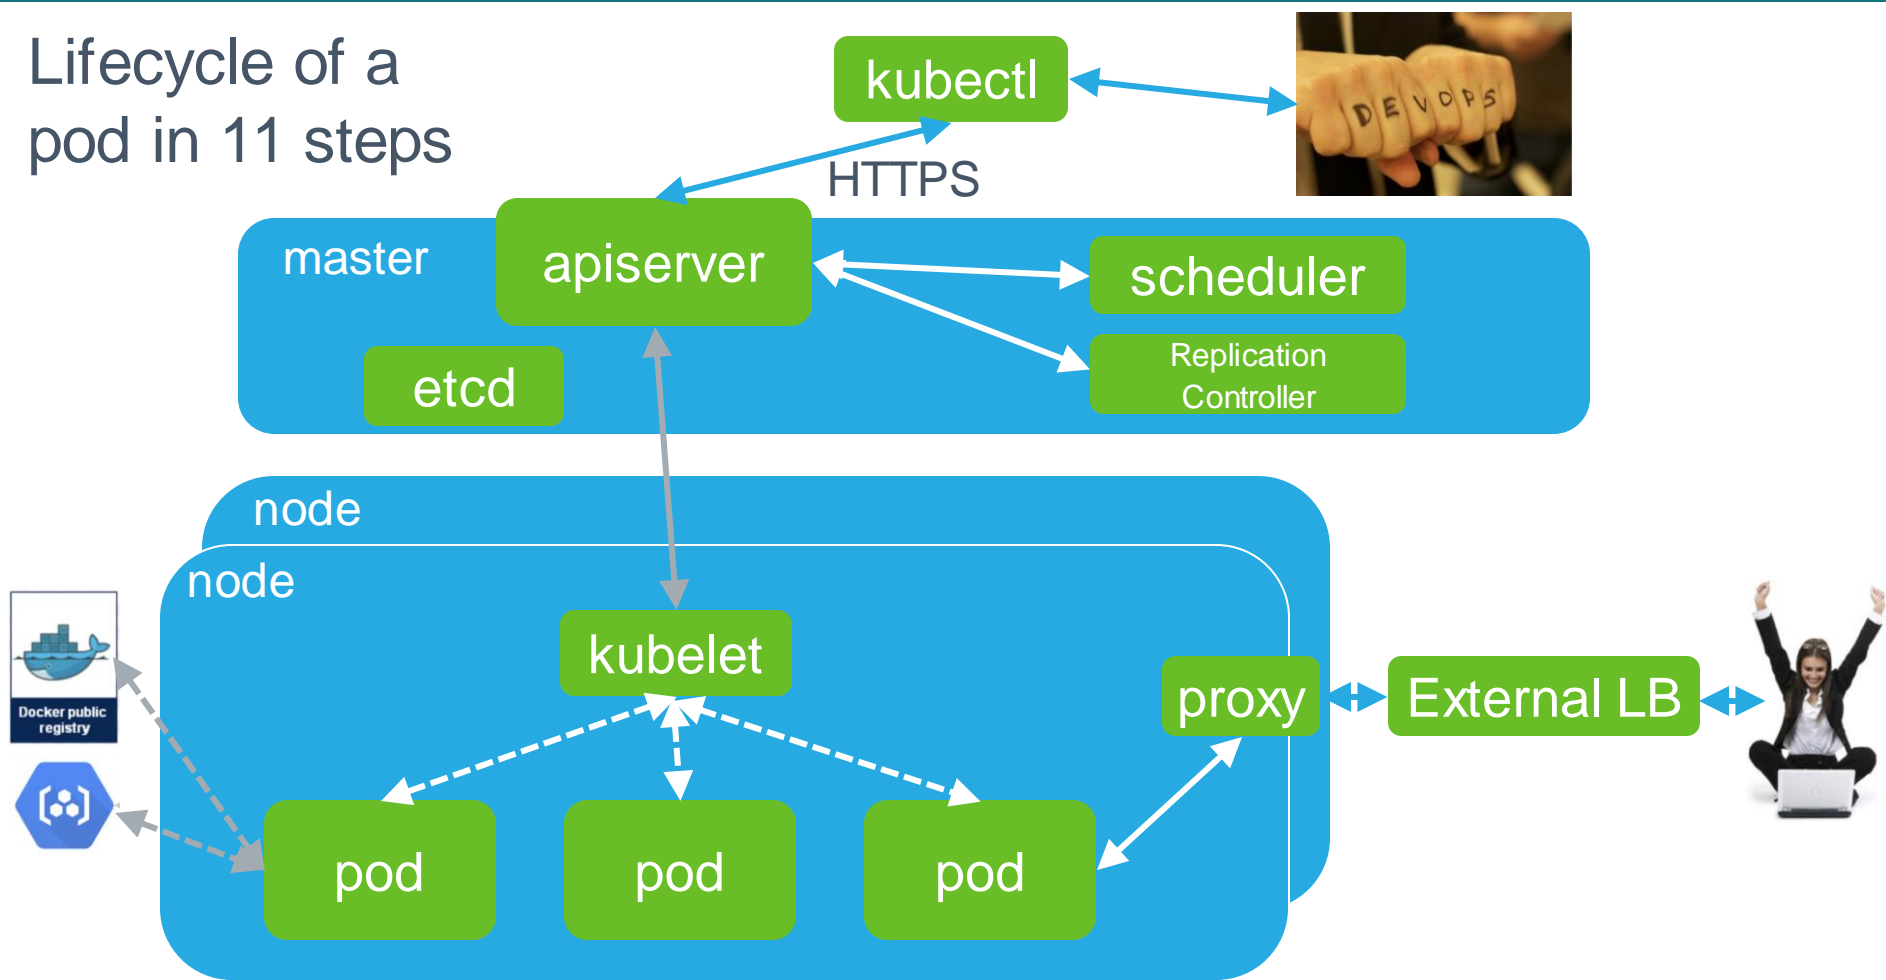
Container runtime

rkt

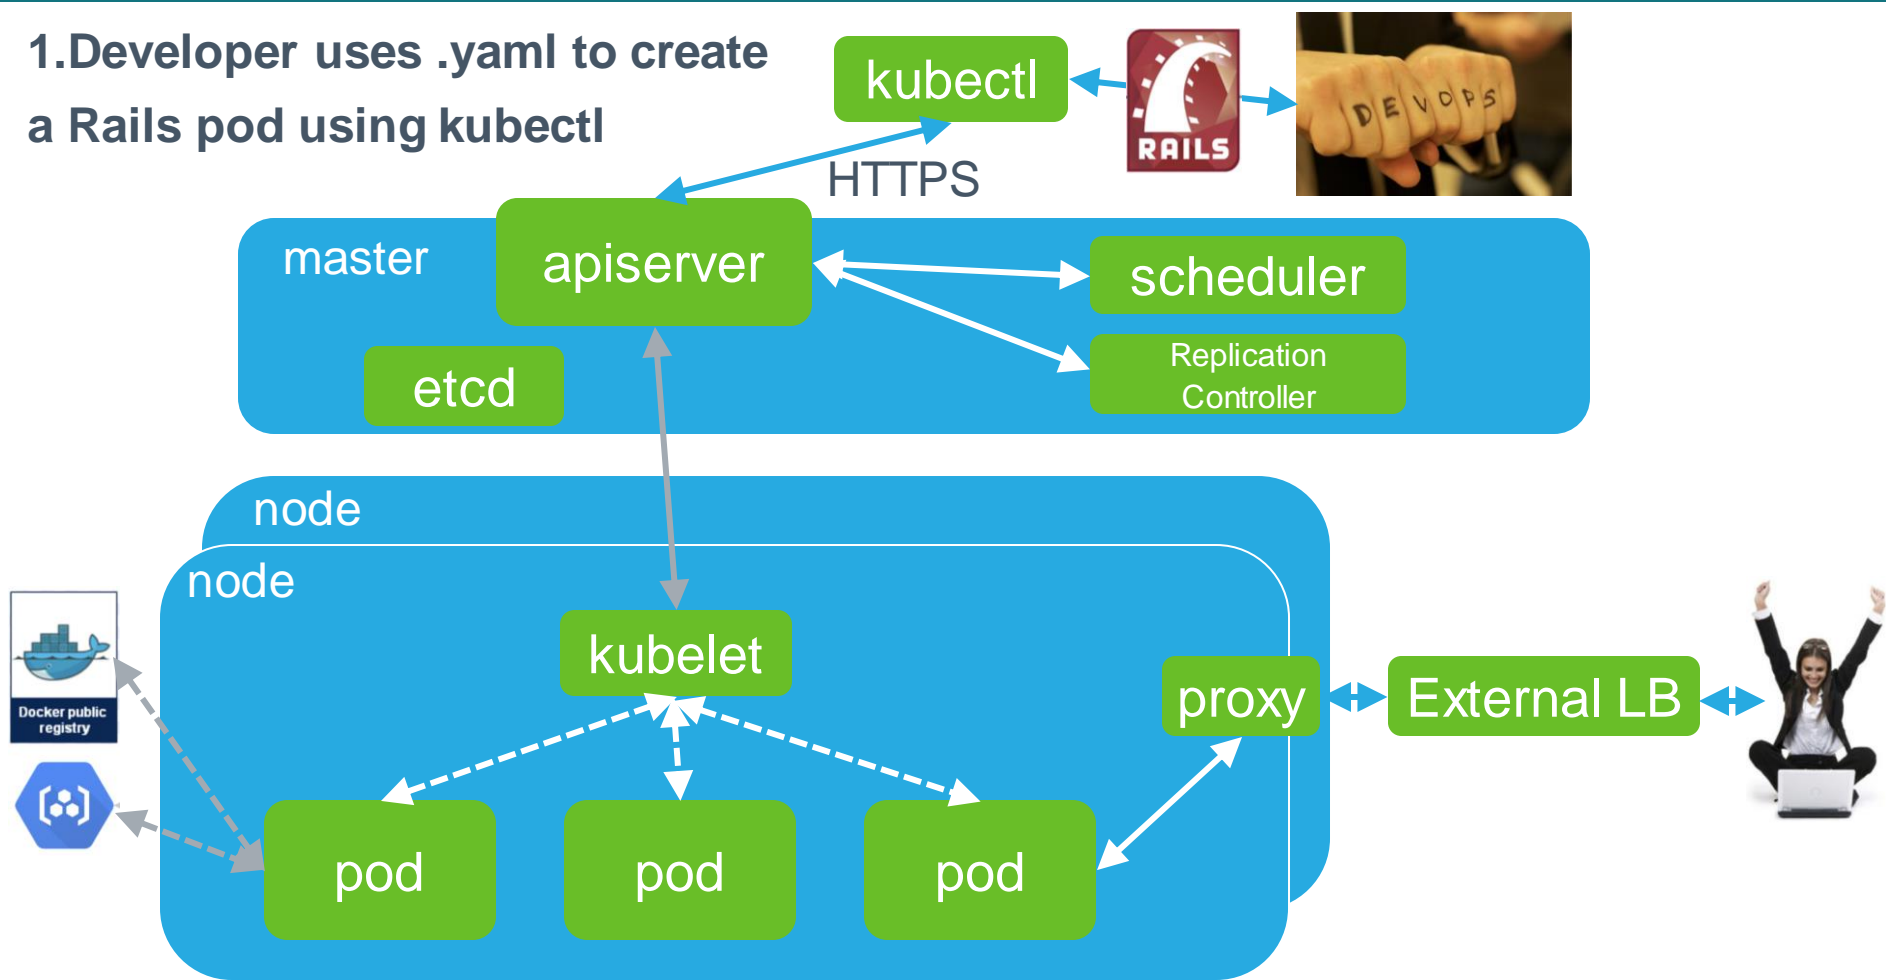


node

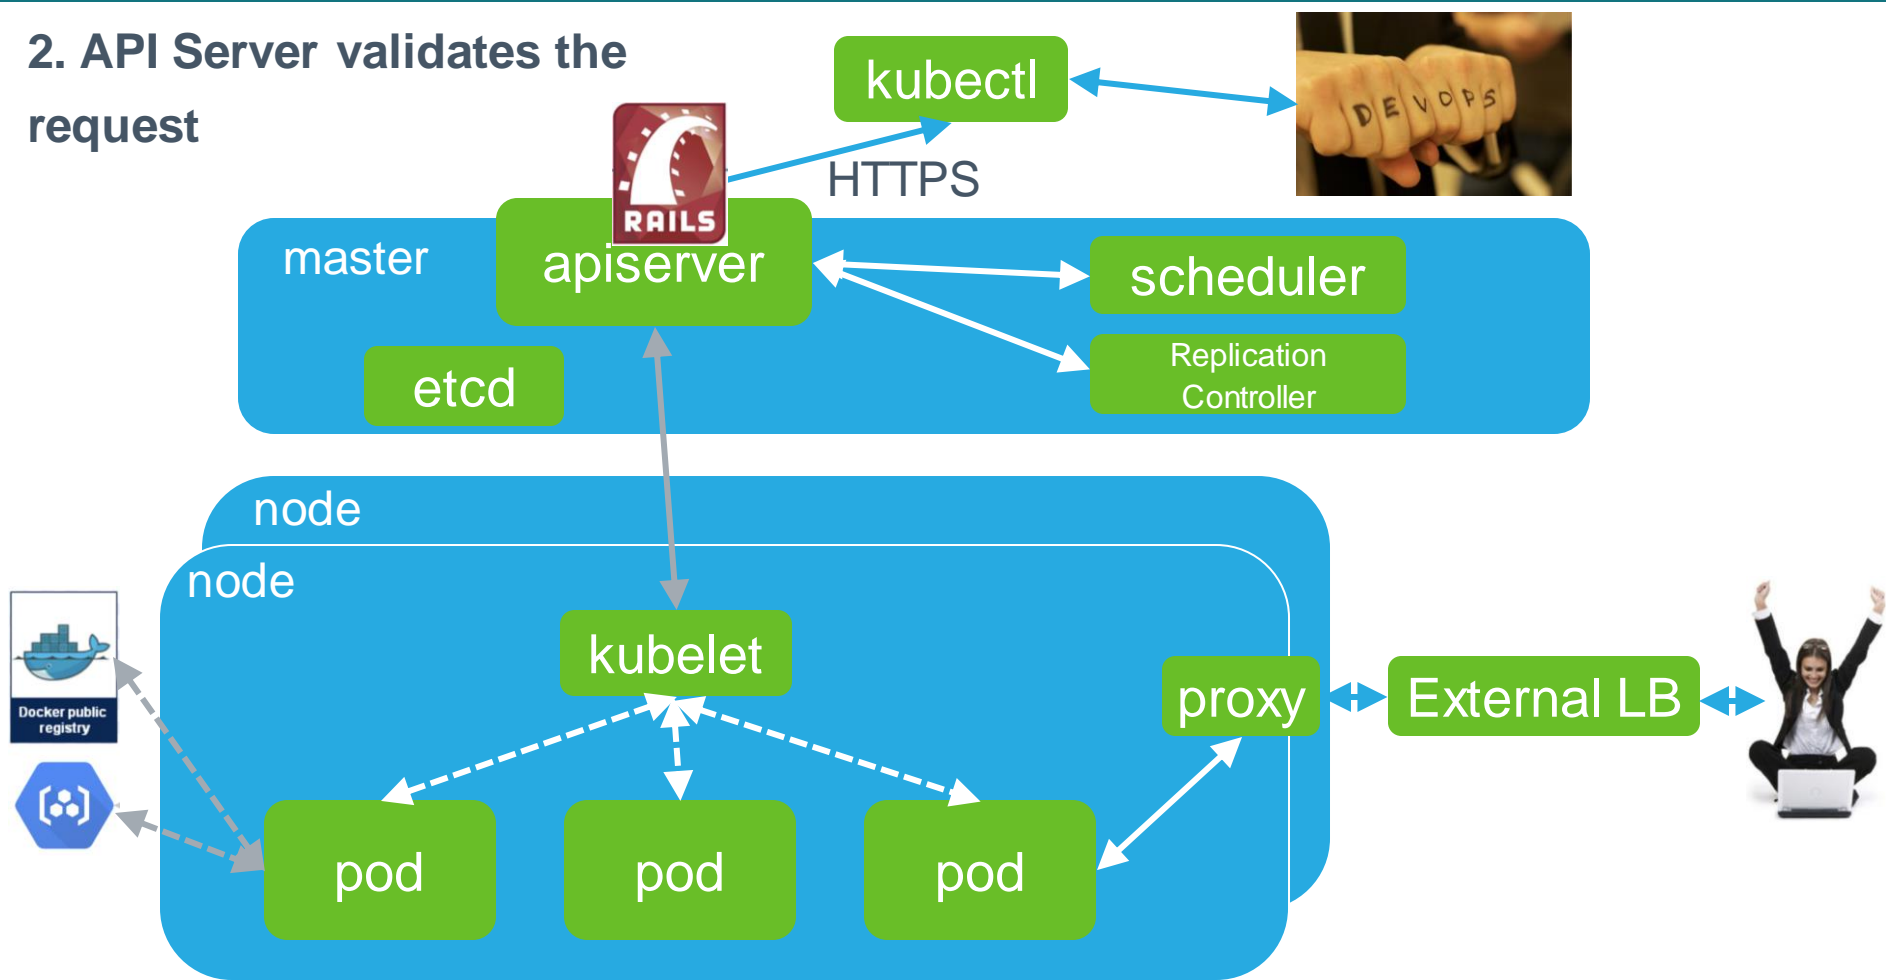
Lifecycle of a pod in 11 steps



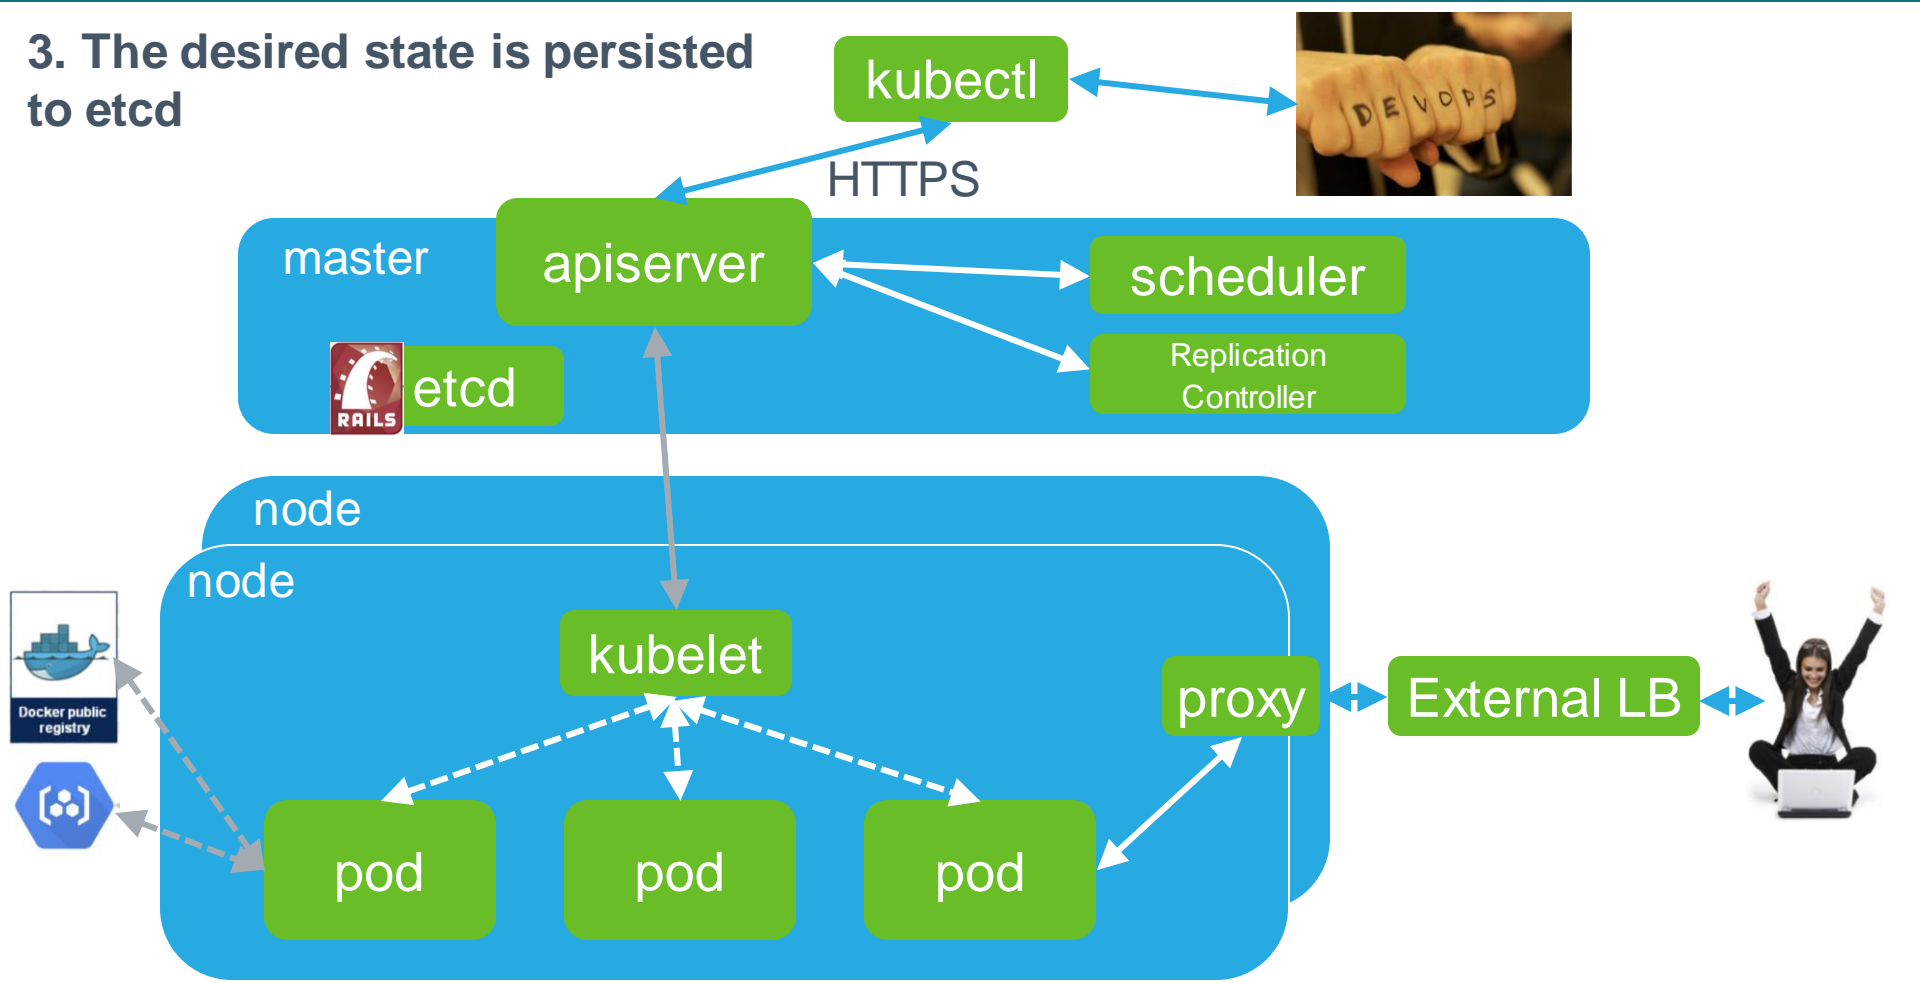
1. Developer uses .yaml to create a Rails pod using kubectl



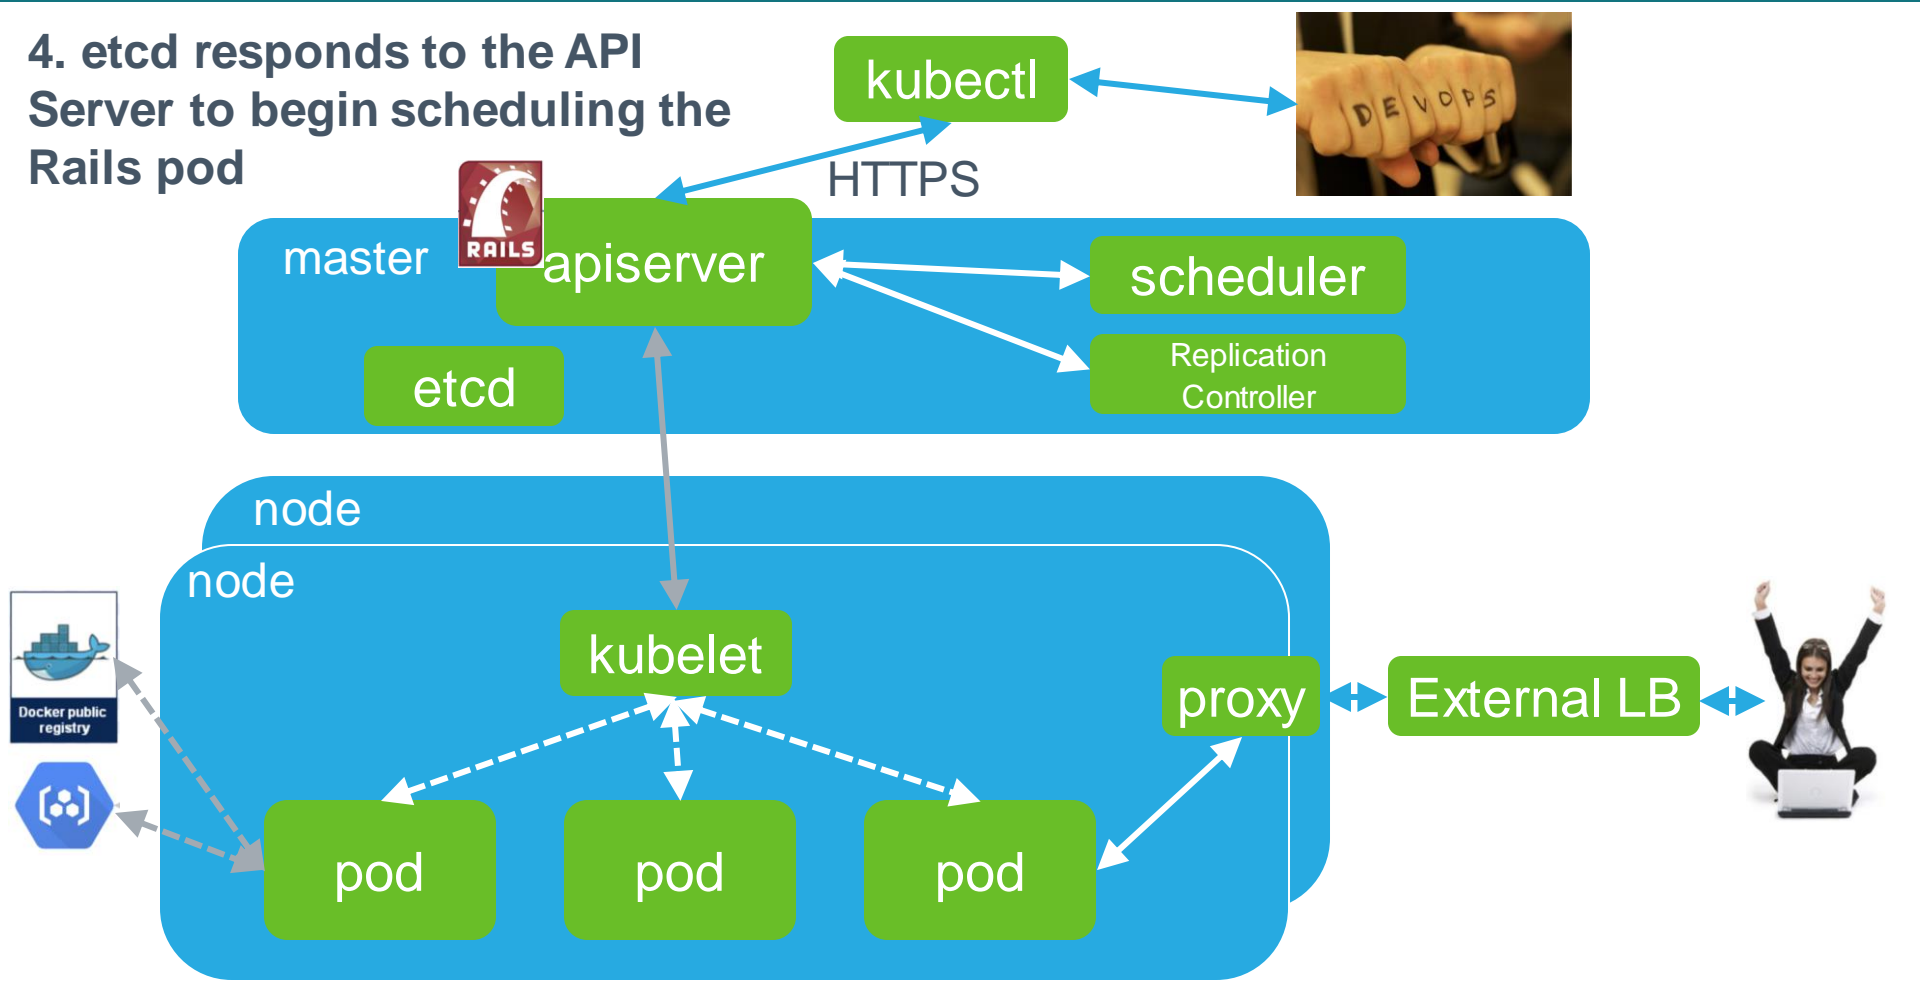
2. API Server validates the request



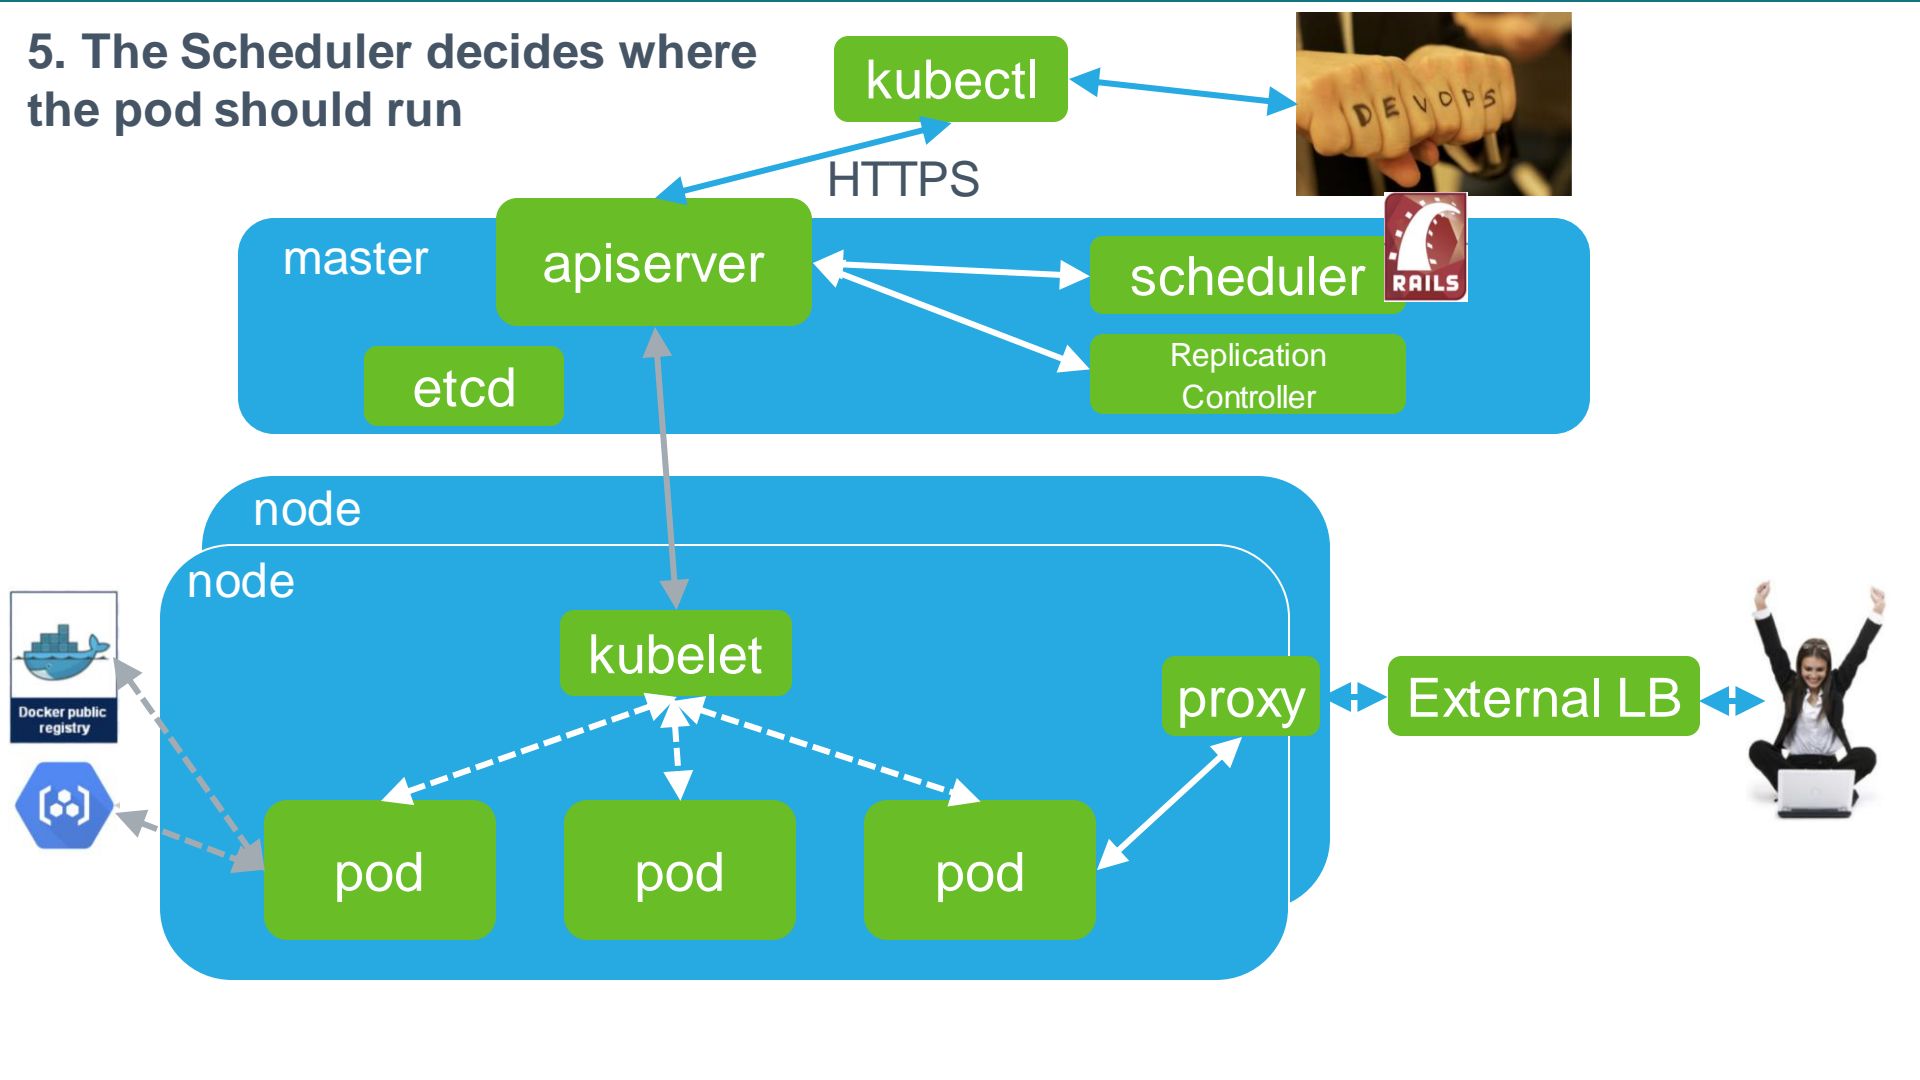
3. The desired state is persisted to etcd




4. etcd responds to the API Server to begin scheduling the Rails pod



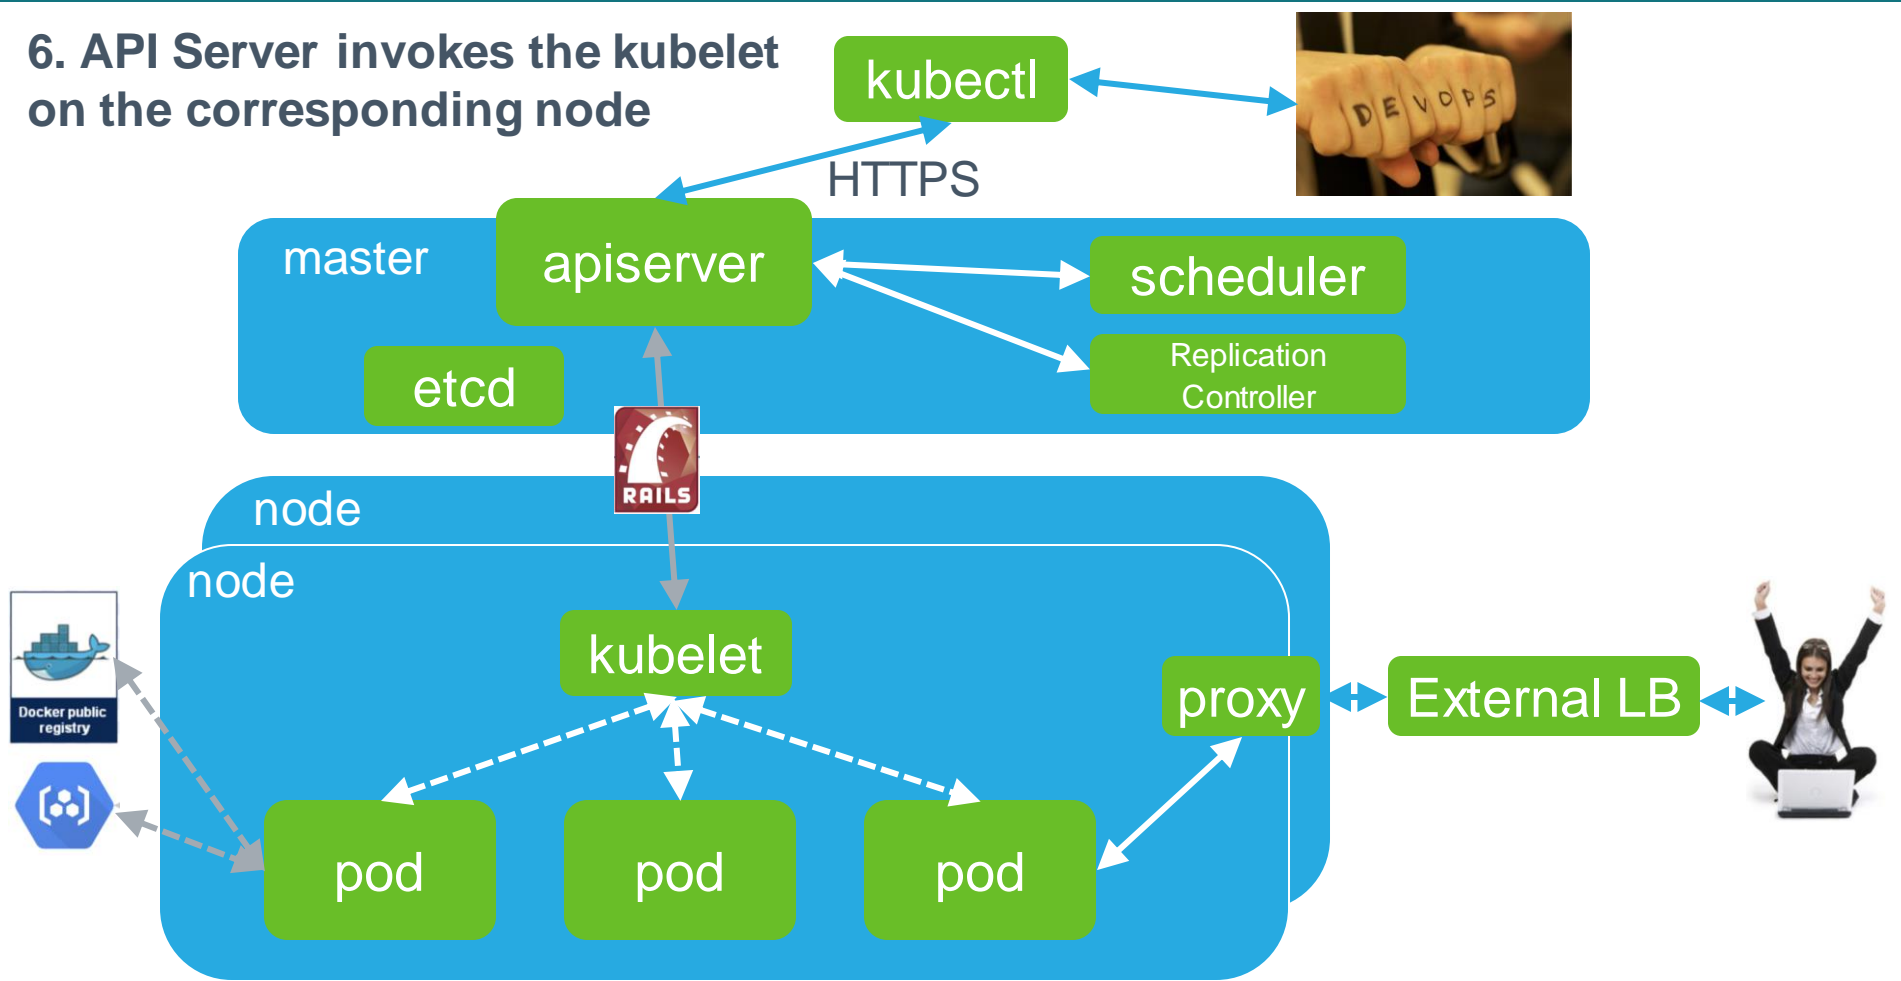
5. The Scheduler decides where the pod should run



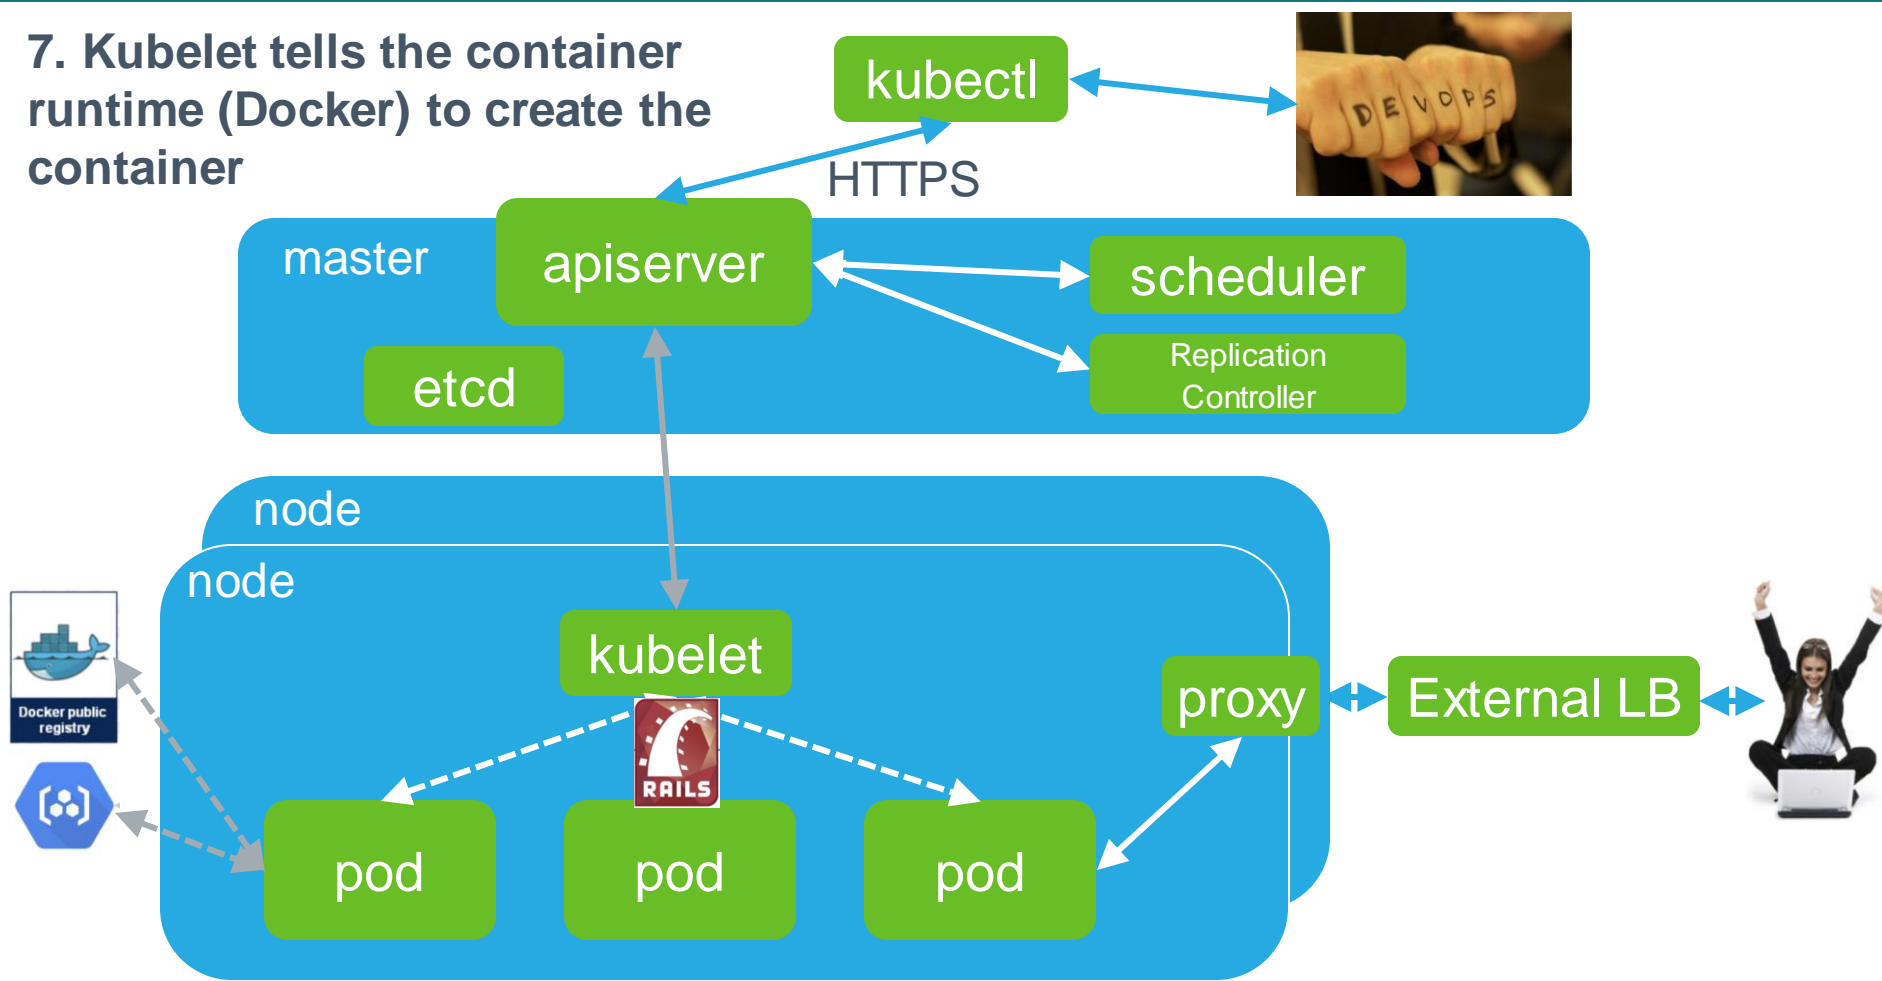
5. The Scheduler decides where the pod should run



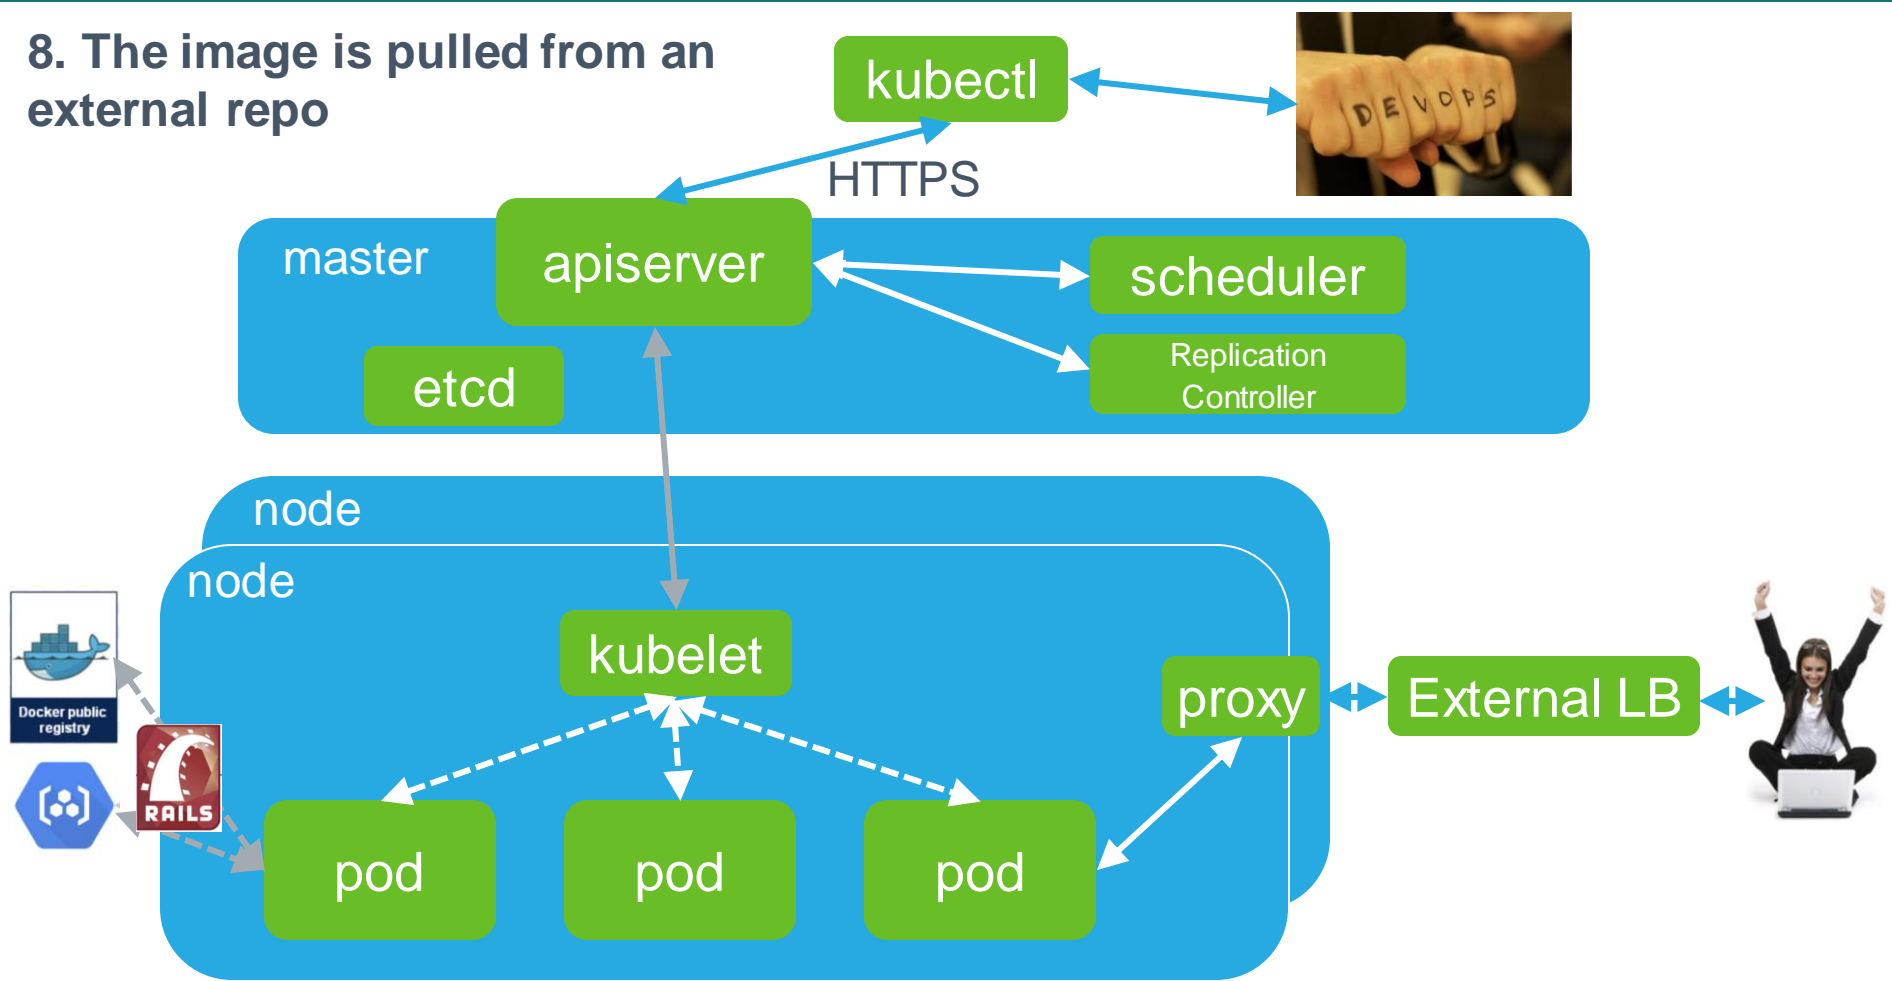
6. API Server invokes the kubelet on the corresponding node



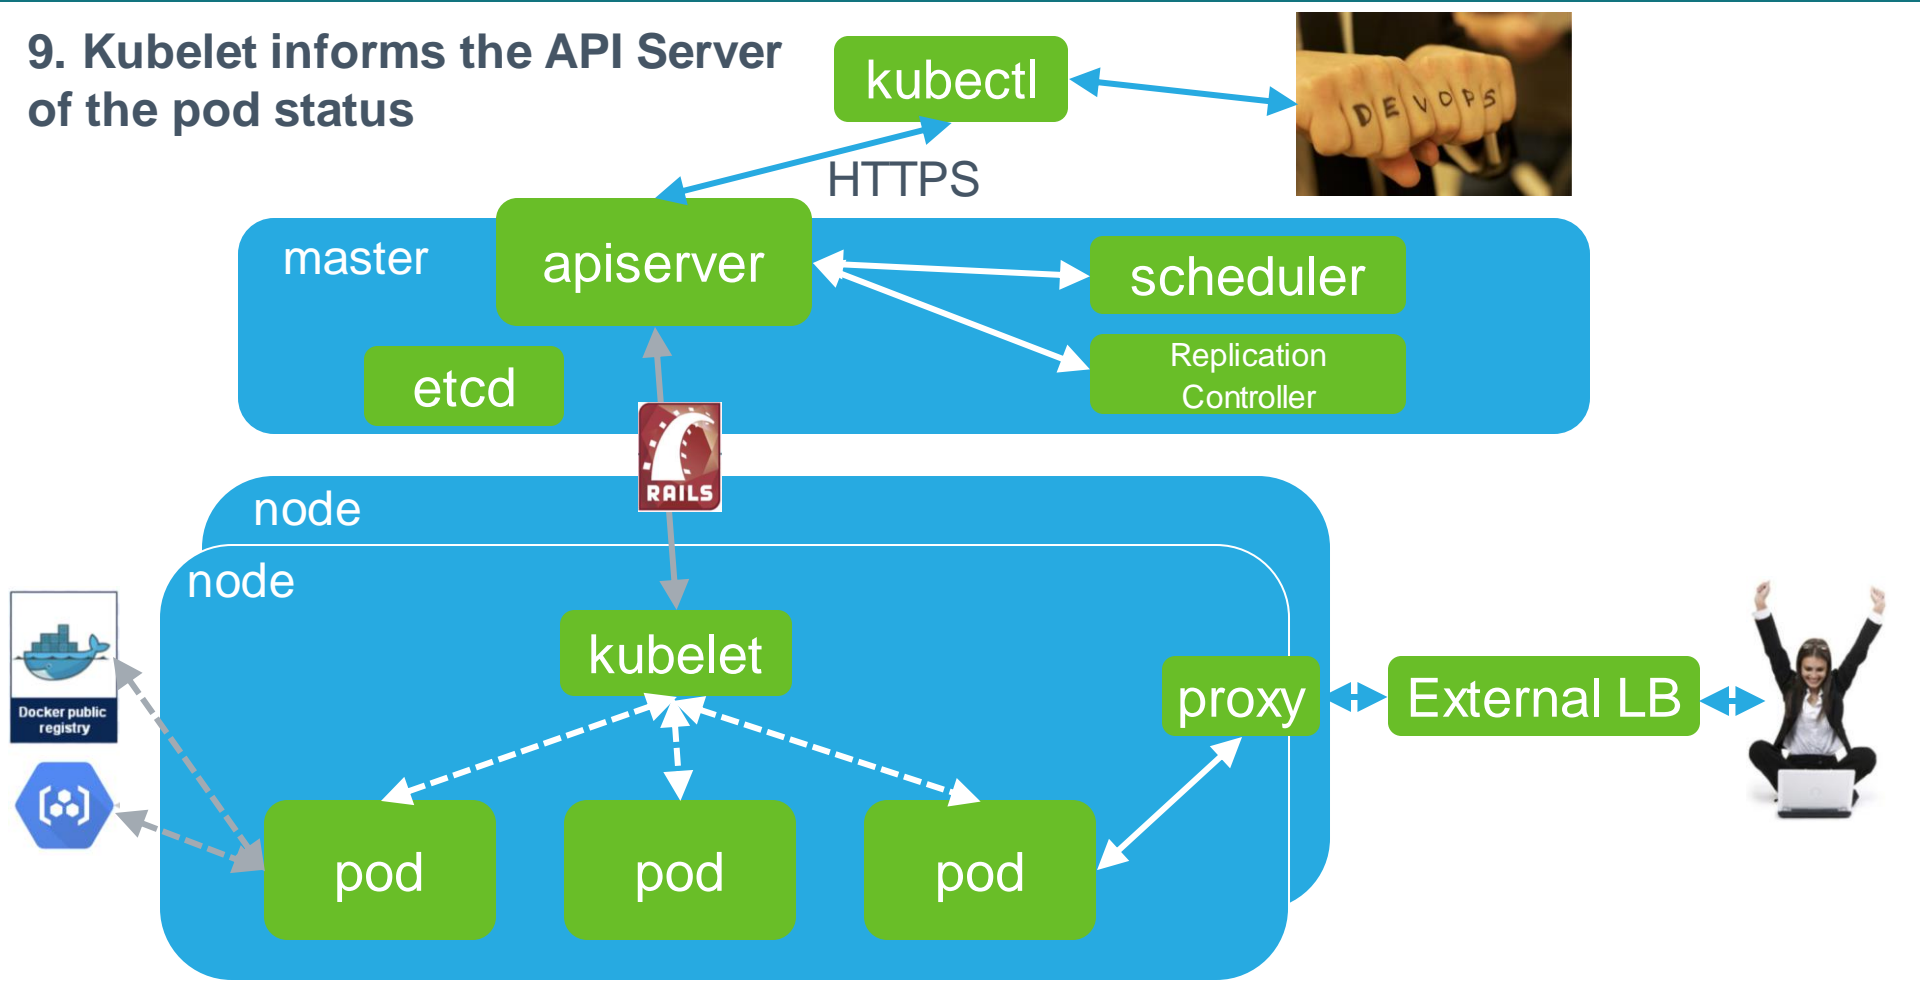
7. Kubelet tells the container runtime (Docker) to create the container



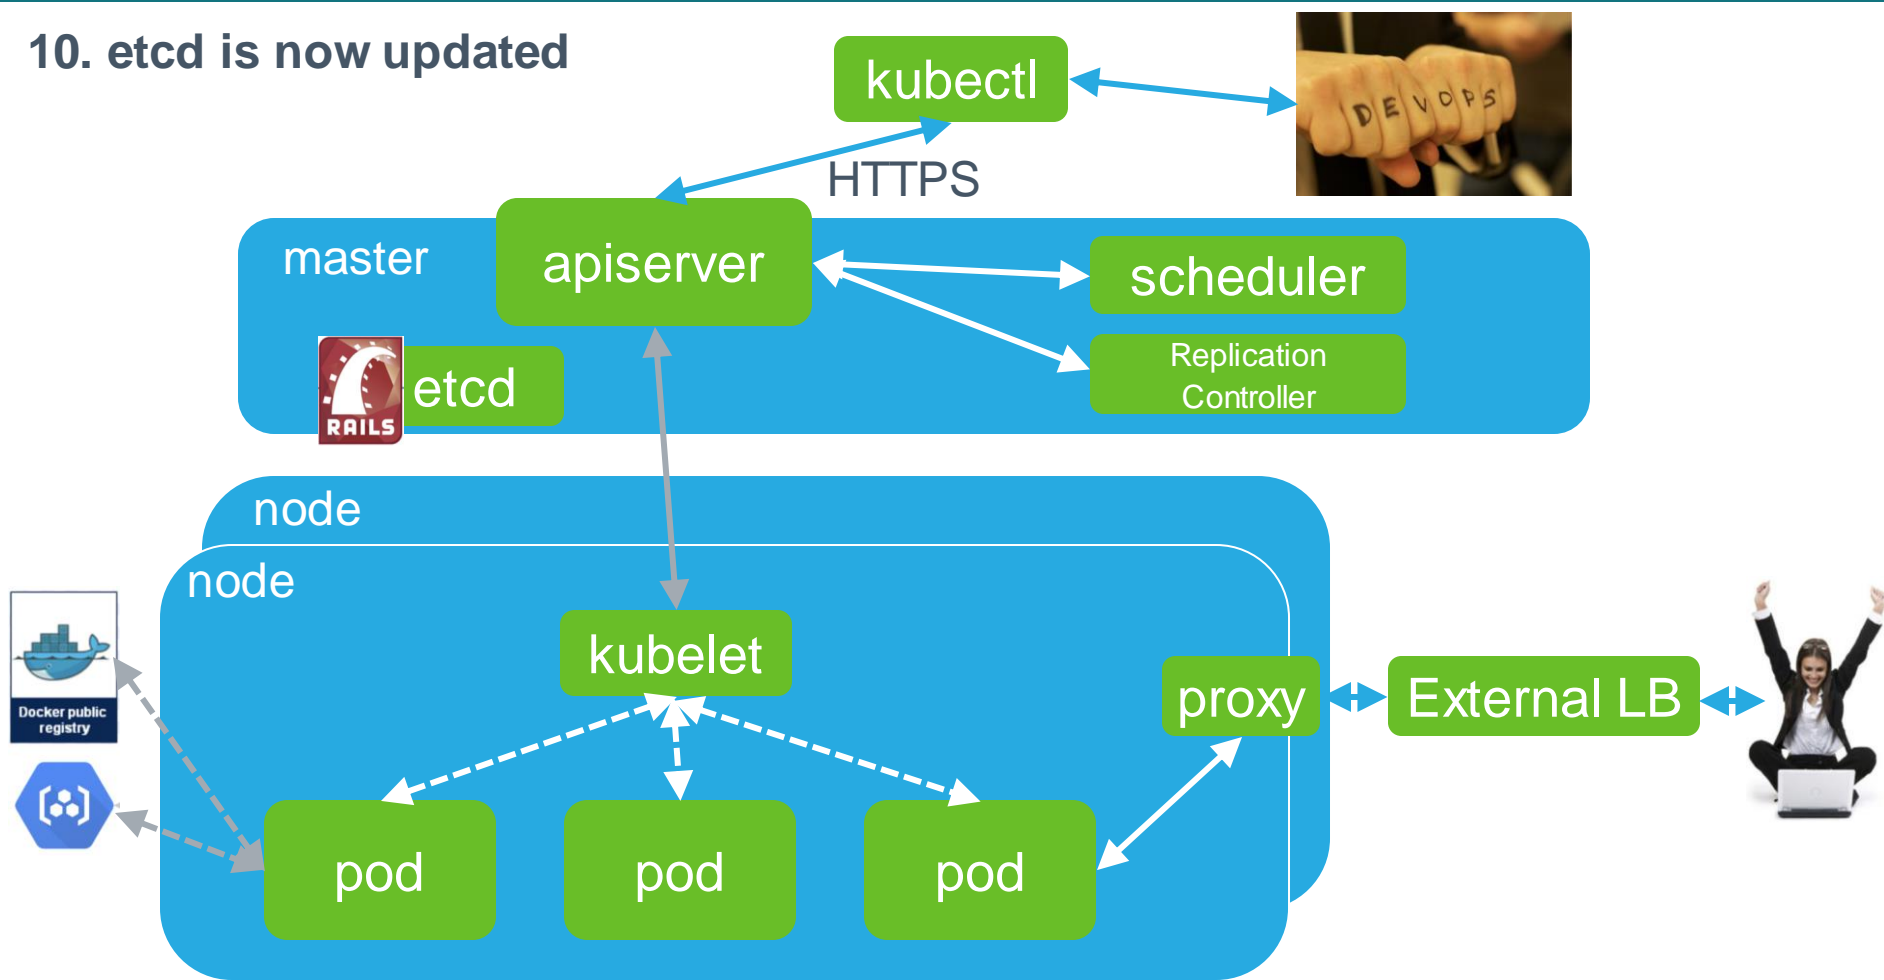
8. The image is pulled from an external repo



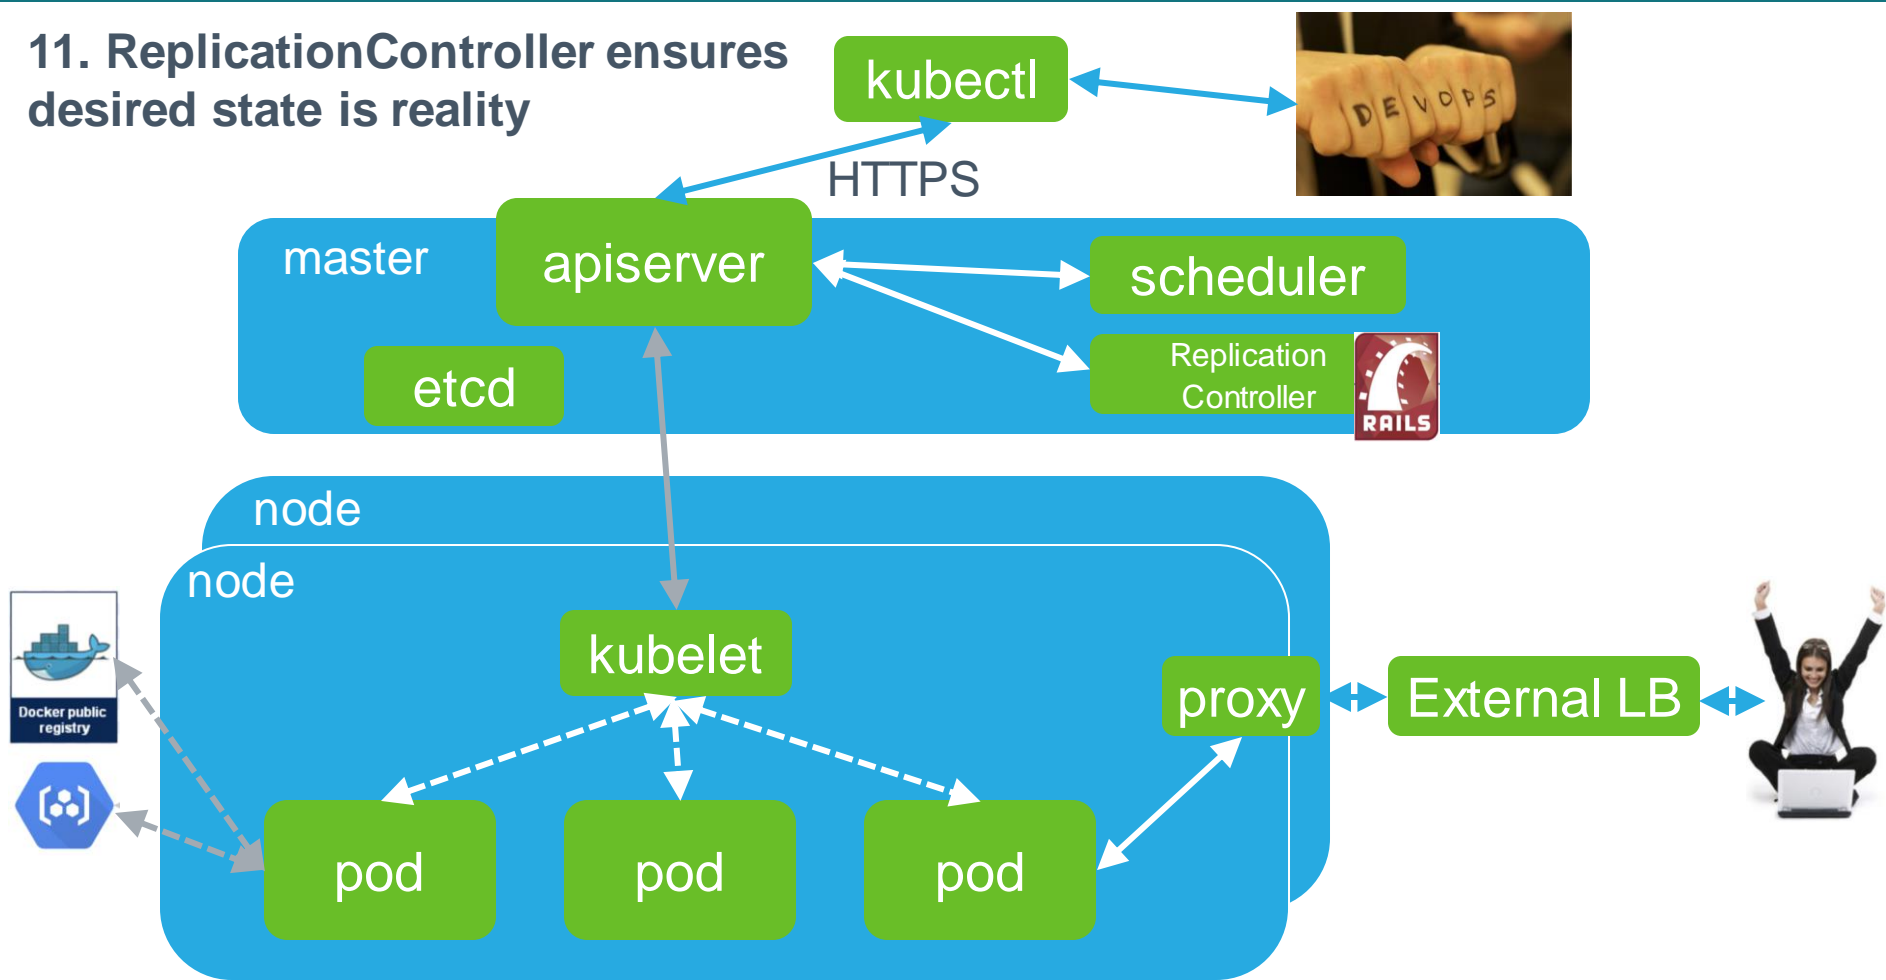
9. Kubelet informs the API Server of the pod status



10. etcd is now updated



11. ReplicationController ensures desired state is reality





Lab 3

003-K8S-Cluster-Setup

Lab Goals:

- Launch API in cluster
- Exec to container
- Expose via LoadBalancer
- Using YAML manifests for deployment